

BIG DATA PIPELINE FOR PASSENGER NAME RECORDS DATA

Subject: Big data storage and processing

Authors

1. Vu Trung Nghia - 20173284
2. Le Vu Loi - 20173240
3. Dang Lam San - 20170111

December 24, 2020

- 1 Introduction
- 2 System architecture
- 3 Execution results
 - Batch processing branch
 - Speed processing branch
- 4 Evaluation
 - KAFKA
 - HDFS
 - SPARK
- 5 Reference

- 1 Introduction
- 2 System architecture
- 3 Execution results
 - Batch processing branch
 - Speed processing branch
- 4 Evaluation
 - KAFKA
 - HDFS
 - SPARK
- 5 Reference



Problem formulation

- Passengers records often come in the form of time series data. On average, there are approximately 3.7 million bookings are performed each day on the targeted systems and therefore the traditional system is no longer sufficient.
- In this project, we built a distributed system and using docker to simulate. Our system has the following constraints:
 - Sufficient for storing and processing big data
 - Fault tolerant handling
 - Easy to expand
 - User can have two views: batch processing and real-time processing. In batch view, user can using SQL api and python to query and visualize data in HDFS, and in real-time view we simple logs the number of received records every 10 seconds.



Record specification

ArrivalTime – local time of arrival

BusinessLeisure – if the trip is for business or leisure

CabinCategory – cabin class

CreationDate – PNR creation date (Julian day)

CurrencyCode – 3-letter currency code of payment

DepartureTime – local time of departure

Destination – IATA code of arrival airport

OfficeIdCountry – country code of office placing the reser

Origin – IATA code of departure airport

TotalAmount – total reservation cost

nPAX – number of passengers


```
{
  "ID": 1188,
  "ArrivalTime": "1453042816",
  "BusinessLeisure": "B",
  "CabinCategory": "40",
  "CreationDate": "2457373",
  "CurrencyCode": "nan",
  "DepartureTime": "1452892672",
  "Destination": "TRD",
  "OfficeIdCountry": "NO",
  "Origin": "ALC",
  "TotalAmount": "nan",
  "nPAX": "1"
}
```



- 1 Introduction
- 2 System architecture**
- 3 Execution results
 - Batch processing branch
 - Speed processing branch
- 4 Evaluation
 - KAFKA
 - HDFS
 - SPARK
- 5 Reference



System architecture



System architecture

- PNR's API produces records to a topic in kafka cluster.
- Batch Processing program and Real-time Processing program subscribe to this topic, these programs use spark cluster 1 and spark cluster 3 as its computation resource.
- Batch Processing program processes data and stores to HDFS. Then user can use Interactive Dashboard which is a notebook to read data from HDFS and do some statistics. This notebook use spark cluster 2 as its computation resource.
- Real-time Processing program processes data and writes to another kafka topic, then the JavaApp subscribes to this topic to read data and write to Graphite database. Finally, Real-time Dashboard reads data from Graphite and visualizes.



HDFS Cluster

```
namenode:
  image: bde2020/hadoop-namenode:2.0.0-hadoop2.7.4-java8
  container_name: namenode
  environment:
    - CLUSTER_NAME=test
  env_file:
    - ./hadoop.env
  ports:
    - 8020:8020
    - 50070:50070
datanode-1:
  image: bde2020/hadoop-datanode:2.0.0-hadoop2.7.4-java8
  container_name: datanode-1
  environment:
    SERVICE_PRECONDITION: "namenode:50070"
  env_file:
    - ./hadoop.env
  ports:
    - 50075:50075
datanode-2:
  image: bde2020/hadoop-datanode:2.0.0-hadoop2.7.4-java8
  container_name: datanode-2
  environment:
    SERVICE_PRECONDITION: "namenode:50070"
  env_file:
    - ./hadoop.env
  ports:
    - 50076:50075
```



Kafka Cluster

```

zookeeper:
  image: zookeeper:3.4.10
  container_name: zookeeper
  environment:
    ZOO_MY_ID: 1
    ZOO_SERVERS: server.1=0.0.0.0:2888:3888
    ZOO_TICK_TIME: 15000
  ports:
    - 2181:2181

kafka-broker-1:
  image: wurstmeister/kafka:2.12-2.2.1
  container_name: kafka-broker-1
  depends_on:
    - zookeeper
  ports:
    - "9092:9092"
  environment:
    - KAFKA_ZOOKEEPER_CONNECT=zookeeper:2181
    - ALLOW_PLAINTEXT_LISTENER=yes
    - KAFKA_ADVERTISED_LISTENERS=INSIDE://kafka-broker-1:9093,OUTSIDE://localhost:9092
    - KAFKA_LISTENER_SECURITY_PROTOCOL_MAP=INSIDE:PLAINTEXT,OUTSIDE:PLAINTEXT
    - KAFKA_LISTENERS=INSIDE://kafka-broker-1:9093,OUTSIDE://0.0.0.0:9092
    - KAFKA_INTER_BROKER_LISTENER_NAME=INSIDE

kafka-broker-2:
  image: wurstmeister/kafka:2.12-2.2.1
  container_name: kafka-broker-2
  depends_on:
    - zookeeper
  ports:
    - "9094:9094"
  environment:
    - KAFKA_ZOOKEEPER_CONNECT=zookeeper:2181
    - ALLOW_PLAINTEXT_LISTENER=yes
    - KAFKA_ADVERTISED_LISTENERS=INSIDE://kafka-broker-2:9093,OUTSIDE://localhost:9094
    - KAFKA_LISTENER_SECURITY_PROTOCOL_MAP=INSIDE:PLAINTEXT,OUTSIDE:PLAINTEXT
    - KAFKA_LISTENERS=INSIDE://kafka-broker-2:9093,OUTSIDE://0.0.0.0:9094
    - KAFKA_INTER_BROKER_LISTENER_NAME=INSIDE
  
```



Spark Cluster

```

post-batch-processing-spark-master:
  image: vutrungnghia99/spark-master:spark2.4.1-python3.7-hadoop2.7
  container_name: post-batch-processing-spark-master
  ports:
    - "8083:8080"
    - "7078:7077"
  environment:
    - INIT_DAEMON_STEP=setup_spark
post-batch-processing-spark-worker-1:
  image: vutrungnghia99/spark-worker:spark2.4.1-python3.7-hadoop2.7
  container_name: post-batch-processing-spark-worker-1
  depends_on:
    - post-batch-processing-spark-master
  environment:
    - "SPARK_MASTER=spark://post-batch-processing-spark-master:7077"
    - "SPARK_WORKER_CORES=1"
    - "SPARK_WORKER_MEMORY=1G"
    - "SPARK_DRIVER_MEMORY=128m"
    - "SPARK_EXECUTOR_MEMORY=256m"
post-batch-processing-spark-worker-2:
  image: vutrungnghia99/spark-worker:spark2.4.1-python3.7-hadoop2.7
  container_name: post-batch-processing-spark-worker-2
  depends_on:
    - post-batch-processing-spark-master
  environment:
    - "SPARK_MASTER=spark://post-batch-processing-spark-master:7077"
    - "SPARK_WORKER_CORES=1"
    - "SPARK_WORKER_MEMORY=1G"
    - "SPARK_DRIVER_MEMORY=128m"
    - "SPARK_EXECUTOR_MEMORY=256m"

```



System manager - Graphite - Grafana

```
##### manager #####
```

```
system-manager:
```

```
  image: vutrungnghia99/system-manager:spark2.4.1-python3.7-hadoop2.7-kafka2.7.0
```

```
  container_name: system-manager
```

```
  ports:
```

```
    - "8888:8888"
```

```
  volumes:
```

```
    - $PWD/src:/home/jovyan/work
```

```
  environment:
```

```
    - JUPYTER_TOKEN=admin
```

```
##### Graphite and Grafana #####
```

```
graphite:
```

```
  image: vutrungnghia99/graphite:1.1.7-6
```

```
  container_name: graphite
```

```
  ports:
```

```
    - "80:80"
```

```
    - "2003:2003"
```

```
    - "2004:2004"
```

```
grafana:
```

```
  image: grafana/grafana:latest
```

```
  container_name: grafana
```

```
  ports:
```

```
    - "3000:3000"
```



- 1 Introduction
- 2 System architecture
- 3 Execution results
 - Batch processing branch
 - Speed processing branch
- 4 Evaluation
 - KAFKA
 - HDFS
 - SPARK
- 5 Reference



- 1 Introduction
- 2 System architecture
- 3 Execution results
 - Batch processing branch
 - Speed processing branch
- 4 Evaluation
 - KAFKA
 - HDFS
 - SPARK
- 5 Reference



List of containers [16]

IMAGE

```

vutrungnghia99/spark-worker:spark2.4.1-python3.7-hadoop2.7
vutrungnghia99/spark-worker:spark2.4.1-python3.7-hadoop2.7
vutrungnghia99/spark-worker:spark2.4.1-python3.7-hadoop2.7
wurstmeister/kafka:2.12-2.2.1
wurstmeister/kafka:2.12-2.2.1
vutrungnghia99/spark-worker:spark2.4.1-python3.7-hadoop2.7
vutrungnghia99/spark-master:spark2.4.1-python3.7-hadoop2.7
zookeeper:3.4.10
vutrungnghia99/system-manager:spark2.4.1-python3.7-hadoop2.7-kafka2.7.0
vutrungnghia99/spark-master:spark2.4.1-python3.7-hadoop2.7
bde2020/hadoop-datanode:2.0.0-hadoop2.7.4-java8
grafana/grafana:latest
vutrungnghia99/spark-master:spark2.4.1-python3.7-hadoop2.7
bde2020/hadoop-namenode:2.0.0-hadoop2.7.4-java8
vutrungnghia99/graphite:1.1.7-6
bde2020/hadoop-datanode:2.0.0-hadoop2.7.4-java8

```

NAMES

```

post-batch-processing-spark-worker-2
post-batch-processing-spark-worker-1
speed-processing-spark-worker-1
kafka-broker-1
kafka-broker-2
pre-batch-processing-spark-worker-1
speed-processing-spark-master
zookeeper
system-manager
post-batch-processing-spark-master
datanode-2
grafana
pre-batch-processing-spark-master
namenode
graphite
datanode-1

```

Spark cluster 1



Spark Master at spark://7a6c33b85408:7077

URL: spark://7a6c33b85408:7077

Alive Workers: 1

Cores in use: 1 Total, 0 Used

Memory in use: 1024.0 MB Total, 0.0 B Used

Applications: 0 Running, 0 Completed

Drivers: 0 Running, 0 Completed

Status: ALIVE

Workers (1)

Worker Id	Address
worker-20201223081126-172.18.0.12-34753	172.18.0.12:34753

Running Applications (0)

Application ID	Name	Cores	Memory per Executor
----------------	------	-------	---------------------

Completed Applications (0)

Application ID	Name	Cores	Memory per Executor
----------------	------	-------	---------------------



BatchProcessing program

```
def get_categorical(x, m):
    if str(x) == 'nan':
        return 0.0
    else:
        v = m['mapping'][str(x)]
        return (v - m['statistic']['mean']) / m['statistic']['std']

def json_to_processed_data(s):
    t = json.loads(s)
    return [
        t['ID'],
        get_continuous(t['ArrivalTime'], mapping_and_statistic['ArrivalTime']),
        get_categorical(t['BusinessLeisure'], mapping_and_statistic['BusinessLeisure']),
        get_categorical(t['CabinCategory'], mapping_and_statistic['CabinCategory']),
        get_continuous(t['CreationDate'], mapping_and_statistic['CreationDate']),
        get_categorical(t['CurrencyCode'], mapping_and_statistic['CurrencyCode']),
        get_continuous(t['DepartureTime'], mapping_and_statistic['DepartureTime']),
        get_categorical(t['Destination'], mapping_and_statistic['Destination']),
        get_categorical(t['OfficeIdCountry'], mapping_and_statistic['OfficeIdCountry']),
        get_categorical(t['Origin'], mapping_and_statistic['Origin']),
        get_continuous(t['TotalAmount'], mapping_and_statistic['TotalAmount']),
        get_continuous(t['nPAX'], mapping_and_statistic['nPAX']),
        s
    ]
```

```
In [*]: ks = KafkaUtils.createDirectStream(
        ssc, [{'metadata.broker.list': 'kafka-broker-1:9093,kafka-broker-2:9093'}])
        lines = ks.map(lambda x: x[1])

        transform1 = lines.map(lambda tripInfo: json_to_list(tripInfo))
        transform1.foreachRDD(handle_rdd1)

        transform2 = lines.map(lambda tripInfo: json_to_processed_data(tripInfo))
        transform2.foreachRDD(handle_rdd2)

        ssc.start()
        ssc.awaitTermination()
```



Spark cluster 2



Spark Master at spark://3d40737aa3f2:7077

URL: spark://3d40737aa3f2:7077

Alive Workers: 2

Cores in use: 2 Total, 0 Used

Memory in use: 2.0 GB Total, 0.0 B Used

Applications: 0 Running, 0 Completed

Drivers: 0 Running, 0 Completed

Status: ALIVE

Workers (2)

Worker Id	Address
worker-20201223081126-172.18.0.16-38487	172.18.0.16:38487
worker-20201223081126-172.18.0.17-37513	172.18.0.17:37513

Running Applications (0)

Application ID	Name	Cores	Memory per Executor
----------------	------	-------	---------------------


Completed Applications (0)



HDFS

Datanode Information

Datanode usage histogram



In operation

Node	Last contact	Admin State	Capacity	Used	Non DFS Used	Remaining	Blocks	Block pool used	Failed Volumes	Version
c65069a49aa9:50010 (172.18.0.3:50010)	1	In Service	355.6 GB	24 KB	83.51 GB	253.96 GB	0	24 KB (0%)	0	2.7.4
5abe514e7d51:50010 (172.18.0.8:50010)	0	In Service	355.6 GB	24 KB	83.51 GB	253.96 GB	0	24 KB (0%)	0	2.7.4

Decommissioning

Node	Last contact	Under replicated blocks	Blocks with no live replicas	Under Replicated Blocks In files under construction
------	--------------	-------------------------	------------------------------	--

3. Execution results

3.1. Batch processing branch

HDFS

[Hadoop](#) [Overview](#) [Datanodes](#) [Snapshot](#) [Startup Progress](#) [Utilities](#) -

Browse Directory

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
Hadoop, 2017.							



Send 10000 records from local machine to topic "trips"

```
Topic: trips    PartitionCount: 2    ReplicationFactor: 2    Configs: segment.bytes=1073741824
Topic: trips    Partition: 0        Leader: 1002    Replicas: 1002,1001    Isr: 1002,1001
Topic: trips    Partition: 1        Leader: 1001    Replicas: 1001,1002    Isr: 1001,1002
```

```
late event
late event
late event
late event
late event
late event
```

```
100%|
```

```
Sent 10000 records in 23.760411262512207 seconds
```

```
Sending rate: 420.86813605694687 records/s
```



```
ation": "LHR", "OfficeIdCountry": "GB", "Origin": "LHR", "TotalAmount": "nan",  
    "nPAx": "1"}H0vooooooooooooooooooooo{"ID": 9981, "ArrivalTime": "144346  
9312", "BusinessLeisure": "nan", "CabinCategory": "40", "CreationDate": "24571  
98", "CurrencyCode": "nan", "DepartureTime": "1443745408", "Destination": "KL  
", "OfficeIdCountry": "DE", "Origin": "MUC", "TotalAmount": "nan", "nPAx": "1"  
}0e06veeeeeooooooooooooooooooooo{"ID": 9982, "ArrivalTime": "1448134912", "Bu  
sinessLeisure": "nan", "CabinCategory": "40", "CreationDate": "2457292", "Cur  
rencyCode": "nan", "DepartureTime": "1445481088", "Destination": "BCN", "Office  
IdCountry": "SE", "Origin": "ARN", "TotalAmount": "nan", "nPAx": "1"}eo{"ID":  
9983, "ArrivalTime": "1437361280", "BusinessLeisure": "nan", "CabinCategory":  
40", "CreationDate": "2457162", "CurrencyCode": "nan", "DepartureTime": "1436  
333312", "Destination": "MLA", "OfficeIdCountry": "FR", "Origin": "CDG", "Tota  
lAmount": "nan", "nPAx": "1"}H0Hooooooooooooooooooooo{"ID": 9985, "Arr  
ivalTime": "1444902656", "BusinessLeisure": "nan", "CabinCategory": "40", "Cre  
ationDate": "2457162", "CurrencyCode": "nan", "DepartureTime": "1445128256", "  
Destination": "JNB", "OfficeIdCountry": "PL", "Origin": "PRG", "TotalAmount":  
"nan", "nPAx": "1"}H0Hooooooooooooooooooooo{"ID": 9986, "ArrivalTime":  
"1439666384", "BusinessLeisure": "nan", "CabinCategory": "40", "CreationDate":  
"2457148", "CurrencyCode": "SAR", "DepartureTime": "1435167232", "Destination":  
": "KHI", "OfficeIdCountry": "SA", "Origin": "RUH", "TotalAmount": "0.0", "nPA  
X": "1"}bash-4.4#
```



The records have been received and processed at Batch Processing Notebook

```
In [*]: ks = KafkaUtils.createDirectStream(
        ssc, ['trips'], {'metadata.broker.list': 'kafka-broker-1:9093,kafka-broker-2:9093'})
        lines = ks.map(lambda x: x[1])

        transform1 = lines.map(lambda tripInfo: json_to_list(tripInfo))
        transform1.foreachRDD(handle_rdd1)

        transform2 = lines.map(lambda tripInfo: json_to_processed_data(tripInfo))
        transform2.foreachRDD(handle_rdd2)

        ssc.start()
        ssc.awaitTermination()
```

```
Recieved 606 records - transform 1
Recieved 606 records - transform 2
Recieved 5832 records - transform 1
Recieved 5832 records - transform 2
Recieved 3562 records - transform 1
Recieved 3562 records - transform 2
```



3. Execution results

3.1. Batch processing branch

Data has been stored in HDFS

Hadoop

Overview

Datanodes

Datanode Volume Failures


Snapshot

Startup Progress

Utilities

Datanode Information

Datanode usage histogram



In operation

Node	Last contact	Admin State	Capacity	Used	Non DFS Used	Remaining	Blocks	Block pool used	Failed Volumes	Version
c65069a49aa9:50010 (172.18.0.3:50010)	2	In Service	355.6 GB	1.91 MB	84.55 GB	252.92 GB	24	1.91 MB (0%)	0	2.7.4
5abe514e7d51:50010 (172.18.0.8:50010)	1	In Service	355.6 GB	1.91 MB	84.55 GB	252.92 GB	24	1.91 MB (0%)	0	2.7.4



3. Execution results

3.1. Batch processing branch

Data has been stored in HDFS

Hadoop Overview Datanodes Snapshot Startup Progress Utilities ▾

Browse Directory

/trips Go!

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
drwxr-xr-x	jovyan	supergroup	0 B	12/23/2020, 3:29:46 PM	0	0 B	processed_trips.parquet
drwxr-xr-x	jovyan	supergroup	0 B	12/23/2020, 3:29:45 PM	0	0 B	trips.parquet

Hadoop, 2017.

3. Execution results

3.1. Batch processing branch

Data has been stored in HDFS

Hadoop Overview Datanodes Snapshot Startup Progress Utilities -								
Browse Directory								
/trips/trips.parquet								Go!
Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name	
-rw-r--r--	joyvan	supergroup	0 B	12/23/2020, 3:29:45 PM	3	128 MB	_SUCCESS	
-rw-r--r--	joyvan	supergroup	88.21 KB	12/23/2020, 3:29:41 PM	3	128 MB	part-00000-19bb1bf0-592d-492d-9595-8487c8c2bcc7-c000.snappy.parquet	
-rw-r--r--	joyvan	supergroup	30.03 KB	12/23/2020, 3:29:45 PM	3	128 MB	part-00000-29e54c4e-d45c-4c86-8dbf-5d6dc4cc17d-c000.snappy.parquet	
-rw-r--r--	joyvan	supergroup	87.09 KB	12/23/2020, 3:29:33 PM	3	128 MB	part-00000-4d646859-ab52-4db8-ad0a-aaec01a80740-c000.snappy.parquet	
-rw-r--r--	joyvan	supergroup	91.86 KB	12/23/2020, 3:29:38 PM	3	128 MB	part-00000-5ad763d6-a1ff-4f85-993c-4c36eb4580f6-c000.snappy.parquet	
-rw-r--r--	joyvan	supergroup	90.76 KB	12/23/2020, 3:29:36 PM	3	128 MB	part-00000-74231608-9c1f-4f52-9d4f-e1ea64568ebb-c000.snappy.parquet	
-rw-r--r--	joyvan	supergroup	42.25 KB	12/23/2020, 3:29:30 PM	3	128 MB	part-00000-c90b3fa1-d391-4476-93b1-ed7d511ed249-c000.snappy.parquet	
-rw-r--r--	joyvan	supergroup	81.64 KB	12/23/2020, 3:29:41 PM	3	128 MB	part-00001-19bb1bf0-592d-492d-9595-8487c8c2bcc7-c000.snappy.parquet	
-rw-r--r--	joyvan	supergroup	29.43 KB	12/23/2020, 3:29:45 PM	3	128 MB	part-00001-29e54c4e-d45c-4c86-8dbf-5d6dc4cc17d-c000.snappy.parquet	
-rw-r--r--	joyvan	supergroup	91.71 KB	12/23/2020, 3:29:33 PM	3	128 MB	part-00001-4d646859-ab52-4db8-ad0a-aaec01a80740-c000.snappy.parquet	
-rw-r--r--	joyvan	supergroup	88.01 KB	12/23/2020, 3:29:39 PM	3	128 MB	part-00001-5ad763d6-a1ff-4f85-993c-4c36eb4580f6-c000.snappy.parquet	



Read data from HDFS and visualize

Read data from parquet file "trips.parquet" in hdfs

```
df = ss.read.parquet("hdfs://namenode:9000/trips/trips.parquet")
```

```
print(f"Number of records: {df.count()}")
df = df.sort('ArrivalTime')
```


Number of records: 10000

Data Mining

Distribution of trips over time

```
arrivalTime = df.select('ArrivalTime').toPandas()['ArrivalTime'].astype('int64')
sns.distplot(arrivalTime)
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f27d875bf90>



- 1 Introduction
- 2 System architecture
- 3 Execution results
 - Batch processing branch
 - Speed processing branch
- 4 Evaluation
 - KAFKA
 - HDFS
 - SPARK
- 5 Reference



Kafka topic => kafka topic

- Program that consumes data from topic "trips" and produces to topic "real-time-statistic"

```
ss = SparkSession.Builder() \
    .appName("SparkBatchStreamingKafka") \
    .master("spark://speed-processing-spark-master:7077") \
    .config("spark.jars", ".:/spark-streaming-kafka-0-10-assembly_2.11-2.4.1.jar,./kafka-clients-0.10.1.0.jar,./spark-sql-kafka-0-10_2.11-2.4.1.jar") \
    .config("spark.sql.warehouse.dir", "hdfs://namenode:9000/") \
    .getOrCreate()

df = ss \
    .readStream \
    .format("kafka") \
    .option("kafka.bootstrap.servers", "kafka-broker-1:9093,kafka-broker-2:9093,kafka-broker-3:9093") \
    .option("partition.assignment.strategy", "none") \
    .option("subscribe", "trips") \
    .load()

import random

def transform_window(s):
    """
    s = Row(start=datetime.datetime(2020, 12, 21, 17, 9, 30), end=datetime.datetime(2020, 12, 21, 17, 9, 40))
    """
    return str(int(s.end.timestamp()))

def transform_count(s):
    """
    s = 941
    """
    return str(s)

udf_transform_window = udf(transform_window)
udf_transform_count = udf(transform_count)

query = df.withWatermark("timestamp", "15 seconds") \
    .groupBy(window("timestamp", "5 seconds", "5 seconds")) \
    .count() \
    .withColumn("count", udf_transform_count("count")) \
    .withColumn("window", udf_transform_window("window")) \
    .withColumn("value", sf.concat(sf.col('window'),sf.lit('_'), sf.col('count')))) \
    .writeStream \
    .format("kafka") \
    .option("kafka.bootstrap.servers", "kafka-broker-1:9093,kafka-broker-2:9093,kafka-broker-3:9093") \
    .option("topic", "real-time-statistic") \
    .option("checkpointLocation", "/tmp/checkpoint") \
    .outputMode("append") \
    .option("truncate", False) \
    .start()

querv.awaitTermination()
```

Kafka topic => graphite

- Program that consumes topic "real-time-statistic" and write to graphite

```

@ localhost:8888/terminal/
jupyter
Logout

group.id = bigdata
heartbeat.interval.ms = 3000
interceptor.classes = null
key.deserializer = class org.apache.kafka.common.serialization.StringDeserializer
max.partition.fetch.bytes = 2640276
max.poll.interval.ms = 300000
max.poll.records = 500
metadata.max.age.ms = 300000
metric.reporters = []
metrics.num.samples = 2
metrics.recording.level = INFO
metrics.sample.window.ms = 30000
partition.assignment.strategy = [class org.apache.kafka.clients.consumer.RangeAssignor]
receive.buffer.bytes = 65536
reconnect.backoff.ms = 50
request.timeout.ms = 300000
retry.backoff.ms = 100
sasl.java.config = null
sasl.kerberos.kinit.cmd = /usr/bin/kinit
sasl.kerberos.min.time.before.relogin = 60000
sasl.kerberos.service.name = null
sasl.kerberos.ticket.renew.jitter = 0.05
sasl.kerberos.ticket.renew.window.factor = 0.8
sasl.mechanism = GSSAPI
security.protocol = PLAINTEXT
send.buffer.bytes = 131072
session.timeout.ms = 180000
ssl.cipher.suites = null
ssl.enabled.protocols = [TLSv1.2, TLSv1.1, TLSv1]
ssl.endpoint.identification.algorithm = null
ssl.key.password = null
ssl.keymanager.algorithm = SunX509
ssl.keystore.location = null
ssl.keystore.password = null
ssl.keystore.type = JKS
ssl.protocol = TLS
ssl.provider = null
ssl.secure.random.implementation = null
ssl.trustmanager.algorithm = PKIX
ssl.truststore.location = null
ssl.truststore.password = null
ssl.truststore.type = JKS
value.deserializer = class org.apache.kafka.common.serialization.StringDeserializer


00:26:17.711 [main] INFO o.a.kafka.common.utils.AppInfoParser - Kafka version: 0.10.2.0
00:26:17.711 [main] INFO o.a.kafka.common.utils.AppInfoParser - Kafka commitId: 976d3baddc0cf421
00:26:17.761 [main] INFO o.a.k.c.c.i.AbstractCoordinator - Discovered coordinator kafka-broker-1:9093 [id: 2147482646 rack
null] for group bigdata
00:26:17.764 [main] INFO o.a.k.c.c.i.ConsumerCoordinator - Revoking previously assigned partitions [] for group bigdata
00:26:17.765 [main] INFO o.a.k.c.c.i.AbstractCoordinator - (Re-)joining group bigdata
00:26:17.837 [main] INFO o.a.k.c.c.i.AbstractCoordinator - Successfully joined group bigdata with generation 1
00:26:17.838 [main] INFO o.a.k.c.c.i.ConsumerCoordinator - Setting newly assigned partitions [real-time-statistic-1, real-
time-statistic-0] for group bigdata

```

3. Execution results

3.2. Speed processing branch

Real-time data in graphite



3. Execution results

3.2. Speed processing branch

Real-time data in grafana



- 1 Introduction
- 2 System architecture
- 3 Execution results
 - Batch processing branch
 - Speed processing branch
- 4 Evaluation**
 - KAFKA
 - HDFS
 - SPARK
- 5 Reference



- 1 Introduction
- 2 System architecture
- 3 Execution results
 - Batch processing branch
 - Speed processing branch
- 4 Evaluation
 - KAFKA
 - HDFS
 - SPARK
- 5 Reference



Stress testing

	2 brokers			1 brokers		
Delay (s)	send rate (records/s)	recieved records	lost records	send rate (records/s)	recieved records	lost records
0	4501.5	10000	0	4990.3	10000	0
	5045.5	10000	0	4945.3	10000	0
	5120.8	10000	0	5003.4	9973	27
0.0000001	1525.6	10000	0	1572.5	10000	0
	1621.1	10000	0	1962.5	10000	0
	1551.4	10000	0	1558.5	10000	0
0.000001	1335.4	10000	0	1126.4	10000	0
	1596.8	10000	0	1866.5	10000	0
	1597.2	10000	0	1581.1	10000	0



Fault tolerant

```
(base) vutrungnghia@Lusheeta:~/kafka_2.13-2.6.0$ bin/kafka-topics.sh --describe --topic trips --bootstrap
p-server localhost:9092,localhost:9094
Topic: trips      PartitionCount: 2      ReplicationFactor: 2      Configs: segment.bytes=1073741824
      Topic: trips Partition: 0      Leader: 1002      Replicas: 1002,1001      Isr: 1002,1001
      Topic: trips Partition: 1      Leader: 1002      Replicas: 1001,1002      Isr: 1002,1001
(base) vutrungnghia@Lusheeta:~/kafka_2.13-2.6.0$ bin/kafka-topics.sh --describe --topic trips --bootstrap
p-server localhost:9092,localhost:9094
Topic: trips      PartitionCount: 2      ReplicationFactor: 2      Configs: segment.bytes=1073741824
      Topic: trips Partition: 0      Leader: 1002      Replicas: 1002,1001      Isr: 1002
      Topic: trips Partition: 1      Leader: 1002      Replicas: 1001,1002      Isr: 1002
(base) vutrungnghia@Lusheeta:~/kafka_2.13-2.6.0$ █
```

Read data from parquet file "trips.parquet" in hdfs

```
df = ss.read.parquet("hdfs://namenode:9000/trips/trips.parquet")
```

```
print(f"Number of records: {df.count()}")
df = df.sort('ArrivalTime')
```

Number of records: 10000



- 1 Introduction
- 2 System architecture
- 3 Execution results
 - Batch processing branch
 - Speed processing branch
- 4 Evaluation
 - KAFKA
 - **HDFS**
 - SPARK
- 5 Reference




Fault tolerant

Hadoop Overview **Datanodes** Datanode Volume Failures Snapshot Startup Progress Utilities ▾

Datanode Information

Datanode usage histogram



In operation

Node	Last contact	Admin State	Capacity	Used	Non DFS Used	Remaining	Blocks	Block pool used	Failed Volumes	Version
f4a27057f907:50010 (172.18.0.3:50010)	1	In Service	355.6 GB	1.97 MB	88.41 GB	249.06 GB	60	1.97 MB (0%)	0	2.7.4
3e4473579333:50010 (172.18.0.8:50010)	1	In Service	355.6 GB	1.97 MB	88.41 GB	249.06 GB	60	1.97 MB (0%)	0	2.7.4



Fault tolerant

```

20/12/23 17:29:34 INFO FileScanRDD: Reading File path: hdfs://namenode:9000/trips/trips.parquet/part-00001-21
range: 0-166672, partition values: [empty row]
20/12/23 17:29:52 WARN BlockReaderFactory: I/O error constructing remote block reader.
java.net.NoRouteToHostException: No route to host
    at sun.nio.ch.SocketChannelImpl.checkConnect(Native Method)
    at sun.nio.ch.SocketChannelImpl.finishConnect(SocketChannelImpl.java:716)
    at org.apache.hadoop.net.SocketIOWithTimeout.connect(SocketIOWithTimeout.java:206)
    at org.apache.hadoop.net.NetUtils.connect(NetUtils.java:531)
    at org.apache.hadoop.hdfs.DFSClient.newConnectedPeer(DFSClient.java:3436)
    at org.apache.hadoop.hdfs.BlockReaderFactory.nextTcpPeer(BlockReaderFactory.java:777)
    at org.apache.hadoop.hdfs.BlockReaderFactory.getRemoteBlockReaderFromTcp(BlockReaderFactory.java:694)
    at org.apache.hadoop.hdfs.BlockReaderFactory.build(BlockReaderFactory.java:355)
    at org.apache.hadoop.hdfs.DFSInputStream.blockSeekTo(DFSInputStream.java:673)
    at org.apache.hadoop.hdfs.DFSInputStream.readWithStrategy(DFSInputStream.java:882)
    at org.apache.hadoop.hdfs.DFSInputStream.read(DFSInputStream.java:934)
    at org.apache.hadoop.hdfs.DFSInputStream.read(DFSInputStream.java:735)
    at java.io.FilterInputStream.read(FilterInputStream.java:83)
    at org.apache.parquet.io.DelegatingSeekableInputStream.read(DelegatingSeekableInputStream.java:61)
    at org.apache.parquet.bytes.BytesUtils.readIntLittleEndian(BytesUtils.java:80)
    at org.apache.parquet.hadoop.ParquetFileReader.readFooter(ParquetFileReader.java:520)
    at org.apache.parquet.hadoop.ParquetFileReader.readFooter(ParquetFileReader.java:505)
    at org.apache.parquet.hadoop.ParquetFileReader.readFooter(ParquetFileReader.java:499)
    at org.apache.parquet.hadoop.ParquetFileReader.readFooter(ParquetFileReader.java:448)
    at

```



Fault tolerant

```
20/12/23 17:31:00 WARN DFSClient: Failed to connect to /172.18.0.3:50010 for block, add to deadNodes
java.net.NoRouteToHostException: No route to host
    at sun.nio.ch.SocketChannelImpl.checkConnect(Native Method)
    at sun.nio.ch.SocketChannelImpl.finishConnect(SocketChannelImpl.java:716)
    at org.apache.hadoop.net.SocketIOWithTimeout.connect(SocketIOWithTimeout.java:206)
    at org.apache.hadoop.net.NetUtils.connect(NetUtils.java:531)
    at org.apache.hadoop.hdfs.DFSClient.newConnectedPeer(DFSClient.java:3436)
    at org.apache.hadoop.hdfs.BlockReaderFactory.nextTcpPeer(BlockReaderFactory.java:777)
    at org.apache.hadoop.hdfs.BlockReaderFactory.getRemoteBlockReaderFromTcp(BlockReaderFactory.java:777)
    at org.apache.hadoop.hdfs.BlockReaderFactory.build(BlockReaderFactory.java:355)
    at org.apache.hadoop.hdfs.DFSInputStream.blockSeekTo(DFSInputStream.java:673)
    at org.apache.hadoop.hdfs.DFSInputStream.readWithStrategy(DFSInputStream.java:882)
    at org.apache.hadoop.hdfs.DFSInputStream.read(DFSInputStream.java:934)
    at org.apache.hadoop.hdfs.DFSInputStream.read(DFSInputStream.java:735)
    at java.io.FilterInputStream.read(FilterInputStream.java:83)
    at org.apache.parquet.io.DelegatingSeekableInputStream.read(DelegatingSeekableInputStream.java:80)
    at org.apache.parquet.bytes.BytesUtils.readIntLittleEndian(BytesUtils.java:80)
    at org.apache.parquet.hadoop.ParquetFileReader.readFooter(ParquetFileReader.java:520)
    at org.apache.parquet.hadoop.ParquetFileReader.readFooter(ParquetFileReader.java:505)
    at org.apache.parquet.hadoop.ParquetFileReader.readFooter(ParquetFileReader.java:499)
    at org.apache.parquet.hadoop.ParquetFileReader.readFooter(ParquetFileReader.java:448)
    at
```



Fault tolerant

```


    at org.apache.spark.sql.execution.BufferedRowIterator.hasNext(BufferedRowIterator.java:43)
    at org.apache.spark.sql.execution.WholeStageCodegenExec$$anonfun$13$$anon$1.hasNext(WholeStageCode
    at scala.collection.Iterator$$anon$11.hasNext(Iterator.scala:409)
    at org.apache.spark.shuffle.sort.BypassMergeSortShuffleWriter.write(BypassMergeSortShuffleWriter.j
    at org.apache.spark.scheduler.ShuffleMapTask.runTask(ShuffleMapTask.scala:99)
    at org.apache.spark.scheduler.ShuffleMapTask.runTask(ShuffleMapTask.scala:55)
    at org.apache.spark.scheduler.Task.run(Task.scala:121)
    at org.apache.spark.executor.Executor$TaskRunner$$anonfun$10.apply(Executor.scala:403)
    at org.apache.spark.util.Utils$.tryWithSafeFinally(Utils.scala:1360)
    at org.apache.spark.executor.Executor$TaskRunner.run(Executor.scala:409)
    at java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1149)
    at java.util.concurrent.ThreadPoolExecutor$Worker.run(ThreadPoolExecutor.java:624)
    at java.lang.Thread.run(Thread.java:748)
20/12/23 17:31:00 INFO DFSClnt: Successfully connected to /172.18.0.8:50010 for BP-1621182885-172.18.0.7
20/12/23 17:31:00 INFO Executor: Finished task 0.0 in stage 4.0 (TID 5). 1716 bytes result sent to driver
20/12/23 17:31:18 INFO CoarseGrainedExecutorBackend: Got assigned task 7
20/12/23 17:31:18 INFO Executor: Running task 0.0 in stage 5.0 (TID 7)
20/12/23 17:31:18 INFO MapOutputTrackerWorker: Updating epoch to 2 and clearing cache
20/12/23 17:31:18 INFO TorrentBroadcast: Started reading broadcast variable 7
20/12/23 17:31:18 INFO MemoryStore: Block broadcast_7_piece0 stored as bytes in memory (estimated size 3.8
20/12/23 17:31:18 INFO TorrentBroadcast: Reading broadcast variable 7 took 13 ms
20/12/23 17:31:18 INFO MemoryStore: Block broadcast_7 stored as values in memory (estimated size 7.1 KB, f
20/12/23 17:31:18 INFO MapOutputTrackerWorker: Don't have map outputs for shuffle 1, fetching them
20/12/23 17:31:18 INFO MapOutputTrackerWorker: Doing the fetch; tracker endpoint = NettyRpcEndpointRef(spa
20/12/23 17:31:18 INFO MapOutputTrackerWorker: Got the output locations
20/12/23 17:31:18 INFO ShuffleBlockFetcherIterator: Getting 2 non-empty blocks including 1 local blocks an
20/12/23 17:31:18 INFO ShuffleBlockFetcherIterator: Started 1 remote fetches in 1 ms
20/12/23 17:31:18 INFO Executor: Finished task 0.0 in stage 5.0 (TID 7). 1782 bytes result sent to driver

```

Fault tolerant

Datanode Information

Datanode usage histogram




In operation

Node	Last contact	Admin State	Capacity	Used	Non DFS Used	Remaining	Blocks	Block pool used	Failed Volumes	Version
4a27057f907:50010 (172.18.0.3:50010)	470	In Service	355.6 GB	2.38 MB	88.43 GB	249.04 GB	116	2.38 MB (0%)	0	2.7.4
3e4473579333:50010 (172.18.0.8:50010)	1	In Service	355.6 GB	2.98 MB	88.43 GB	249.03 GB	116	2.98 MB (0%)	0	2.7.4

Fault tolerant

Datanode Information

Datanode usage histogram



In operation

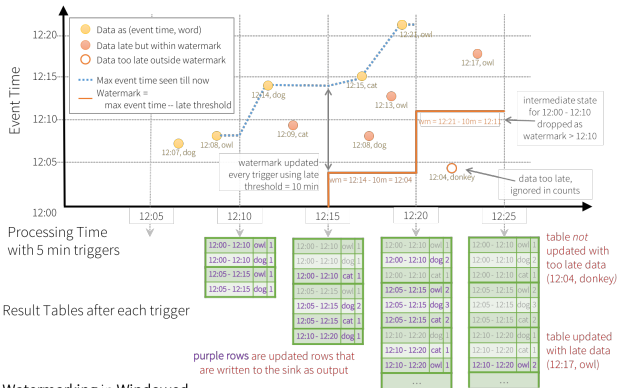
Node	Last contact	Admin State	Capacity	Used	Non DFS Used	Remaining	Blocks	Block pool used	Failed Volumes	Version
3e4473579333:50010 (172.18.0.8:50010)	2	In Service	355.6 GB	2.98 MB	88.44 GB	249.02 GB	116	2.98 MB (0%)	0	2.7.4
14a27057f907:50010 (172.18.0.3:50010)	Thu Dec 24 2020 00:29:21 GMT+0700 (Indochina Time)	Dead	-	-	-	-	-	-	-	-



- 1 Introduction
- 2 System architecture
- 3 Execution results
 - Batch processing branch
 - Speed processing branch
- 4 Evaluation
 - KAFKA
 - HDFS
 - SPARK
- 5 Reference



Late messages handling



(*) source: <https://spark.apache.org>



Late messages handling

```
query = df.withWatermark("timestamp", "20 seconds") \
    .groupBy(window("timestamp", "10 seconds", "10 seconds")) \
    .count() \
```

- To handle late event, Structured Spark supports Window Grouped Aggregation and Watermarking.
- Watermarking excludes event that is too late (for some threshold).



Late messages handling

```

Row(start=datetime.datetime(2020, 12, 22, 4, 22), end=datetime.datetime(2020, 12, 22, 4, 22, 10)) 108
Row(start=datetime.datetime(2020, 12, 22, 4, 22, 10), end=datetime.datetime(2020, 12, 22, 4, 22, 20)) 138
Row(start=datetime.datetime(2020, 12, 22, 4, 22, 20), end=datetime.datetime(2020, 12, 22, 4, 22, 30)) 51
Row(start=datetime.datetime(2020, 12, 22, 4, 22, 30), end=datetime.datetime(2020, 12, 22, 4, 22, 40)) 106
Row(start=datetime.datetime(2020, 12, 22, 4, 22, 40), end=datetime.datetime(2020, 12, 22, 4, 22, 50)) 135
Row(start=datetime.datetime(2020, 12, 22, 4, 22, 50), end=datetime.datetime(2020, 12, 22, 4, 23)) 56
Row(start=datetime.datetime(2020, 12, 22, 4, 23), end=datetime.datetime(2020, 12, 22, 4, 23, 10)) 103
Row(start=datetime.datetime(2020, 12, 22, 4, 23, 10), end=datetime.datetime(2020, 12, 22, 4, 23, 20)) 136
Row(start=datetime.datetime(2020, 12, 22, 4, 23, 20), end=datetime.datetime(2020, 12, 22, 4, 23, 30)) 58
Row(start=datetime.datetime(2020, 12, 22, 4, 23, 30), end=datetime.datetime(2020, 12, 22, 4, 23, 40)) 104
Row(start=datetime.datetime(2020, 12, 22, 4, 23, 40), end=datetime.datetime(2020, 12, 22, 4, 23, 50)) 134
Row(start=datetime.datetime(2020, 12, 22, 4, 23, 50), end=datetime.datetime(2020, 12, 22, 4, 24)) 59

```

Figure: Count the number of records for the last 10 seconds

- Using KafkaProducer to simulate late event.
- After 20s (watermarking threshold), half of the records are assigned as late event.



Reference

- <https://github.com/haiphucnguyen/BigDataDemo>
- <http://www.diva-portal.org/smash/get/diva2:897808/FULLTEXT01.pdf>
- <https://blog.softwaremill.com/7-mistakes-when-using-apache-spark-2017-02-28/>
- <https://spark.apache.org/docs/1.5.2/sql-programming-guide.html>
- <https://spark.apache.org/docs/latest/structured-streaming-programming-guide.html>
- <https://medium.com/dev-genius/an-in-depth-look-at-zookeeper-2017-02-28/>
- <https://spark.apache.org/docs/latest/streaming-programming-guide.html>

